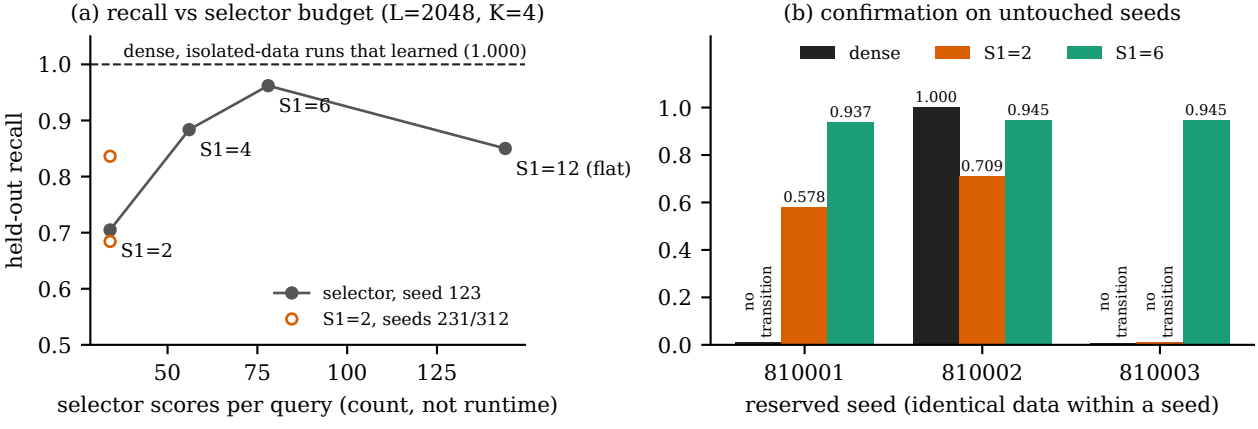
(a) recall vs selector budget (L=2048, K=4)
dense, isolated-data runs that learned (1.000)
held-out recall
1.0
0.9
0.8
0.7
0.6
0.5
S1=6
S1=4
S1=12 (flat)
S1=2
selector, seed 123
S1=2, seeds 231/312
50
75
100
125
selector scores per query (count, not runtime)
(b) confirmation on untouched seeds
dense
S1=2
S1=6
1.0
0.8
0.6
0.4
0.2
0.0
0.937
0.578
1.000
0.709
0.945
0.945
no transition
no transition
no transition
810001
810002
810003
reserved seed (identical data within a seed)